

CUDA-Based Fibonacci Sequence

Goal:

The objective was to:

- Implement the Fibonacci sequence computation in CUDA.
- Compare CUDA's parallel implementation with a sequential CPU implementation.
- Analyze performance, scalability, and gain deeper understanding of GPU parallelism for a simple dependency-based algorithm.

My Approach:

1. Understanding the Problem

The Fibonacci sequence is a recursive problem where each term depends on the previous two terms.

While the formula looks simple, the dependency between terms creates a serial dependency chain that is not naturally parallel.

2. Using LLM to brainstorm parallel approaches:

LLM Prompt I gave to GPT:

"How can I write a CUDA kernel for computing Fibonacci sequence for N numbers? Is it parallelizable?"

GPT's Insight:

- Fibonacci sequence is not naturally parallel for standard computation.
- Two options:
 - Use recursion and memoization on CPU side.
 - Use wavefront parallelism or prefix sum-like techniques to parallelize Fibonacci in CUDA.
 - Or simply compute independent Fibonacci numbers using the closed-form Binet's formula (approximate but parallel).

3. Implemented two versions:

- a. Sequential CPU version (Baseline):
 - i. Simple loop-based code.
 - ii. $O(N)$ time complexity.
- b. CUDA GPU version:
 - i. Each CUDA thread computes one Fibonacci number independently.
 - ii. Used a simple iterative loop inside each thread.
 - iii. Even though threads are independent, each thread computes its own Fibonacci number serially.

4. Code Setup

- a. Used Google Colab for CUDA access (Tesla K80 GPUs).
- b. CUDA kernel written in Fibonacci.ipynb.

CUDA-Based Fibonacci Sequence

- c. Both CPU and GPU versions tested for same inputs (N up to 220).
- d. Verified correctness of results between CPU & GPU versions.

I generated simple output graphs (see README.md and notebook) to visualize execution time differences as problem size grows.

Insights:

- Algorithms like Fibonacci highlight when hardware acceleration makes sense.
- Data dependency limits the gains of parallelism.
- CUDA shines best for tasks where computations across data elements are independent.