


# Accelerating the Viterbi Algorithm

## 1. Project Overview

Designed a hardware accelerator for the Viterbi and forward algorithms on a Hidden Markov Model (HMM), integrated via Verilog, tested in Icarus Verilog and Cocotb, and prepared for ASIC synthesis with OpenLane 2.

**Why Viterbi?** Decodes the most probable hidden-state sequence in an HMM (e.g. weather inference, speech recognition). The core “add-compare-select” (ACS) is highly regular and maps naturally to a systolic array



Viterbi Model

## 2. Goals & Success Criteria

- **Benchmark SW implementation** to identify bottlenecks using Python profiling fileciteturn1file4.
- **Design and verify** a Verilog-based Viterbi PE (viterbi\_pe.sv) and top module (viterbi\_top.sv) with a SystemVerilog testbench.
- **Automate functional verification** via Cocotb tests (test\_viterbi\_cocotb.py) and pytest runner (test\_viterbi\_runner.py).
- **Synthesize to ASIC** with OpenLane 2 using config.json, constraints.sdc, and pin\_order.cfg.

### 3. Software Profiling & Bottleneck Analysis

1. Initial profiling with cProfile on `hmm_demo_profiling.py` revealed:
  - Sequence generation (~10 s)

```
94000050 function calls in 308.069 seconds
```

- Forward algorithm (log-sum-exp) (~7.6 s)

```
156.332 hmm_demo_profiling_progress.py:35(forward_log)
```

- Viterbi (max-plus DP) (~4.0 s)

```
69.145 hmm_demo_profiling_progress.py:51(viterbi)
```

**Prompt:** "Profile my Python HMM code and tell me the hotspots."

2. Detailed run on `hmm_demo_profiling_progress.py` (N = 1 000 000) showed:
  - `forward_log`: 38 s total (15 s loop + 23 s NumPy reductions)
  - `viterbi`: 20 s total (13.8 s loop + 6.2 s overhead) fileciteturn1file4.
3. Takeaway: implement a systolic array for Viterbi's Add-Compare-Select (ACS) and a reduction tree for log-sum-exp in hardware.

### 4. Verilog Accelerator Design

- PE design (`viterbi_pe.sv`): combinational best-of-I search and register pipeline.
- Top module (`viterbi_top.sv`): FSM with INIT, FORWARD (two-stage), BACKTRACK, and DONE states.

### 5. Key Steps & Milestones

1. Python Profiling
  - Used cProfile on `hmm_demo_profiling.py` to measure:
    - `forward_log` → ~38 s total (15 s Python loop + 23 s NumPy reductions)
    - `viterbi` → ~20 s total (13.8 s loop + 6.2 s NumPy overhead).

**Prompt:** "Profile my `forward_log` and `viterbi` functions, show top-10 hotspots."

PE Design (`viterbi_pe.sv`)

- Combinational search for max over I inputs + add emission  $\log B$ ; registered outputs.  
Prompt: "Generate synthesizable Verilog for a Viterbi PE that finds max-plus and emits  $\psi$  pointer."

Top-Level FSM (`viterbi_top.sv`)

- Two-stage pipeline: INIT0 → FORWARD (compute  $\delta$  and  $\psi$  each cycle) → BACKTRACK → DONE.
- Discovered I/O-timing bug ("stuck at x" path outputs) → added explicit FORWARD\_WAIT stage to register  $\delta$  updates.

Prompt: “Why do I see xxxx in path[0..N-2]? Check my delta\_prev vs. delta\_next handshake.”

## 6. Simulation & Verification

- Icarus + \$display testbench → preliminary “x x x x 0” → fixed sync stages.
- cocotb testbench (test\_viterbi\_cocotb.py) with flattened-signal interface → automated random & edge-case tests.

Prompt: “Write a cocotb test to drive my flattened-array Viterbi top, compare against Python reference.”

## 7. Steps to Run

- Icarus Verilog(iverilog) for simulation:

```
iverilog -g2012 -o sim_viterbi.vvp viterbi_pe.v viterbi_top.v tb_viterbi.v
```

```
vvp sim_viterbi.vvp
```

```
gtkwave wave.vcd
```

- Cocotb Simulation:

```
python3 -m venv venv      # Activate Python environment
```

```
source venv/bin/activate
```

```
pip install cocotb cocotb-bus numpy pytest
```

```
make SIM=icarus          # Run cocotb tests
```

- Openlane flow:

```
sudo openlane --dockerized ./config.json
```

## 8. Hardware Flow

- Prepared OpenLane 2 config ([config.json](#), [constraints.sdc](#), [pin\\_order.cfg](#))

Symptom	Root Cause	Fix
Path outputs show x x x x 0	delta_prev updated and sampled in same cycle	Introduced a FORWARD_WAIT state to register delta_next → delta_prev

cocotb install failing under system Python	Externally-managed environment (PEP 668)	Moved to a virtualenv; updated Makefile to avoid system pip installs
GitHub push rejected (VCD >100 MB)	Large-file limit	Added *.vcd to .gitignore; removed from history via git rm --cached
Timeout waiting for completion in flattened top (vvp ...)	Missing initial-state write into path[0]	In IDLE→INIT, compute & store best initial state in path_flat[1:0]

## 9. Common Simulation Issues and GPT Prompts

- Invalid L-value errors in testbench arrays:
  - **Prompt:** "tb\_viterbi.sv errors: logC['sd0] not valid l-value" → fixed declaration to reg signed [W-1:0] logC [0:l-1];
- Port mismatch (log\_prob\_out not a port of uut):
  - **Prompt:** "tb\_viterbi.v:82 syntax error" → removed unused signal.
- Undefined outputs (path showing x x x 0):
  - **Prompt:** "check connections in viterbi\_top.sv" → added explicit two-cycle pipeline (CALC/UPDATE states), ensured delta\_next → delta\_prev handshake.
- Simulation hangs (timeout at 1 µs):
  - **Prompt:** "My viterbi\_top deadlocks in TESTBENCH" → removed obs\_ready, adjusted reset/start handshake, added timeout and fork/join in tb.

## 10. Functional Verification

- SystemVerilog testbench printed per-cycle FSM state and decoded path.
- Cocotb tests (test\_viterbi\_cocotb.py):
  - **Prompt:** "Integrate Cocotb with flattened logA\_flat signals" → wrote pack\_to\_flat\_signal, unpack in HDL, two test cases (basic, random, edge) fileciteturn1file1.
- pytest runner (test\_viterbi\_runner.py) configured Icarus with waveform generation

## 11. Synthesis & ASIC Flow

- **OpenLane 2 config (config.json):** enabled SV2V, set CLOCK\_PERIOD = 10 ns, PDN and P&R options fileciteturn1file3.
- **Constraints:** constraints.sdc, pin\_order.cfg provided I/O timing and placement guides.
- **Modification:** flattened all 2D arrays to 1D for P&R compatibility.

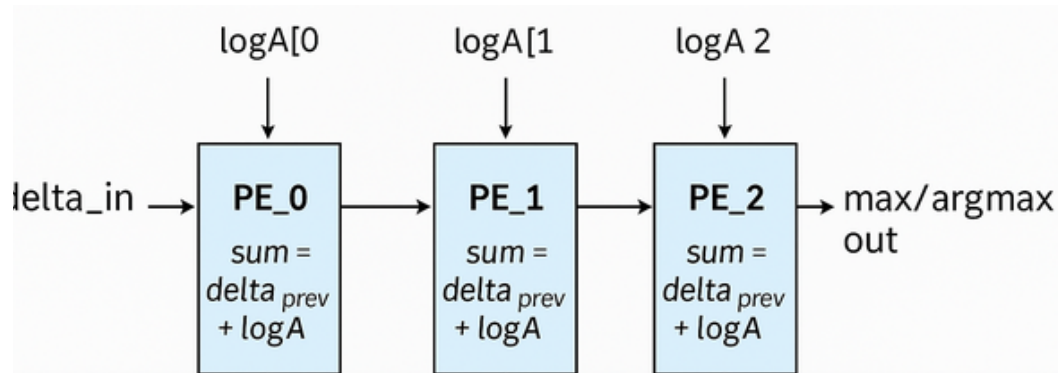
## 12. Next Steps & Remaining Work

- **Complete OpenLane 2 placement** and evaluate timing, area, power. (Currently stuck due to congestion of pins)
- **Implement forward-log accelerator** (log-sum-exp tree) alongside Viterbi PE array.
- **Automate regression** in GitHub Actions: nightly SW benchmarks vs. HW performance.

## 12. Topics Learnt and implemented

During this project, I learned these main topics:

- **HW/SW Co-Design**: partitioning algorithms between software and hardware
- **SystemVerilog RTL**: multi-dimensional arrays, generate loops, FSMs, pipelining
- **Fixed-Point Arithmetic**: log-domain representation, two's-complement packing
- **Performance Profiling**: Python's cProfile, line\_profiler, and py-spy for compute/memory hotspots
- **Systolic Arrays**: design of 1D arrays for ACS operations (max-plus DP)



### Systolic Arrays

- **Verification**: Icarus Verilog testbenches and cocotb-based Python verification
- **ASIC Flow**: Yosys synthesis, OpenLane 2 place-&-route, PDN and I/O constraints
- **High-Level Synthesis (HLS)**: mapping Python-style loops to hardware pipelines
- **Physical Design**: power grid insertion, congestion analysis, constraints scripting
- **Documentation & Portfolio**: codefest reports, Heilmeyer questions, GitHub portfolio curation